

Virtual Memory

Prof. Eun-Seok Ryu (esryu@skku.edu)

Multimedia Computing Systems Laboratory

<http://mcs.l.skku.edu>

Department of Immersive Media Engineering

Department of Computer Education (J.A.)

Sungkyunkwan University (SKKU)

Background

- Code needs to be in memory to execute, but entire program rarely used
 - Error code, unusual routines, large data structures
- Entire program code not needed at same time
- Consider ability to execute partially-loaded program
 - Program no longer constrained by limits of physical memory
 - Each program takes less memory while running -> more programs run at the same time
 - ▶ Increased CPU utilization and throughput with no increase in response time or turnaround time
 - Less I/O needed to load or swap programs into memory -> each user program runs faster

- Q: program 들은 크고 여러개이고
동시 사용하려 할 때, main memory
한계로 다 못 올리면 어떻게?
> 가상메모리


Background (Cont.)

- **Virtual memory** – separation of user logical memory from physical memory
 - Only part of the program needs to be in memory for execution
 - Logical address space can therefore be much larger than physical address space
 - Allows address spaces to be shared by several processes
 - Allows for more efficient process creation
 - More programs running concurrently
 - Less I/O needed to load or swap processes

Background (Cont.)


- **Virtual address space** – logical view of how process is stored in memory
 - Usually start at address 0, contiguous addresses until end of space
 - Meanwhile, physical memory organized in page frames
 - MMU must map logical to physical
- Virtual memory can be implemented via:
 - Demand paging
 - Demand segmentation

Virtual Memory That is Larger Than Physical Memory




Virtual-address Space

- Usually design logical address space for stack to start at Max logical address and grow “down” while heap grows “up”
 - Maximizes address space use
 - Unused address space between the two is hole
 - ▶ No physical memory needed until heap or stack grows to a given new page
- Enables sparse address spaces with holes left for growth, dynamically linked libraries, etc
- System libraries shared via mapping into virtual address space
- Shared memory by mapping pages read-write into virtual address space
- Pages can be shared during fork(), speeding process creation




Shared Library Using Virtual Memory



Demand Paging

- Could bring entire process into memory at load time
- Or bring a page into memory only when it is needed
 - Less I/O needed, no unnecessary I/O
 - Less memory needed
 - Faster response
 - More users
- Similar to paging system with swapping (diagram on right)
- Page is needed \Rightarrow reference to it
 - invalid reference \Rightarrow abort
 - not-in-memory \Rightarrow bring to memory
- **Lazy swapper** – never swaps a page into memory unless page will be needed
 - Swapper that deals with pages is a **pager**



- Demand Paging: 필요한 페이지만 메모리에 올리고, 필요하지 않은 페이지는 디스크에 저장하여 메모리를 절약하는 방법

- Q: Page단위로 swap in/out될 때, 필요한 page가 main memory에 있는지 빠르게 알 수 있는가?

Basic Concepts

- With swapping, pager guesses which pages will be used before swapping out again
- Instead, pager brings in only those pages into memory
- How to determine that set of pages?
 - Need new MMU functionality to implement demand paging
- If pages needed are already **memory resident**
 - No difference from non demand-paging
- If page needed and not memory resident
 - Need to detect and load the page into memory from storage
 - ▶ Without changing program behavior
 - ▶ Without programmer needing to change code

Valid-Invalid Bit


- With each page table entry a valid–invalid bit is associated (**v** \Rightarrow in-memory – **memory resident**, **i** \Rightarrow not-in-memory)
- Initially valid–invalid bit is set to **i** on all entries
- Example of a page table snapshot:

Frame #	valid-invalid bit
	v
	v
	v
	i
...	
	i
	i

page table

- During MMU address translation, if valid–invalid bit in page table entry is **i** \Rightarrow page fault

Page Table When Some Pages Are Not in Main Memory



Page Fault

- If there is a reference to a page, first reference to that page will trap to operating system:

page fault


1. Operating system looks at another table to decide:

- Invalid reference \Rightarrow abort
- Just not in memory

• 본 과정을 뒷페이지
그림으로 설명

2. Find free frame
3. Swap page into frame via scheduled disk operation
4. Reset tables to indicate page now in memory
Set validation bit = **V**
5. Restart the instruction that caused the page fault

Steps in Handling a Page Fault



- Q: page swap in/out이
여기 저기 일어나면 효율이
문제되진 않을까?
> locality of reference

Aspects of Demand Paging

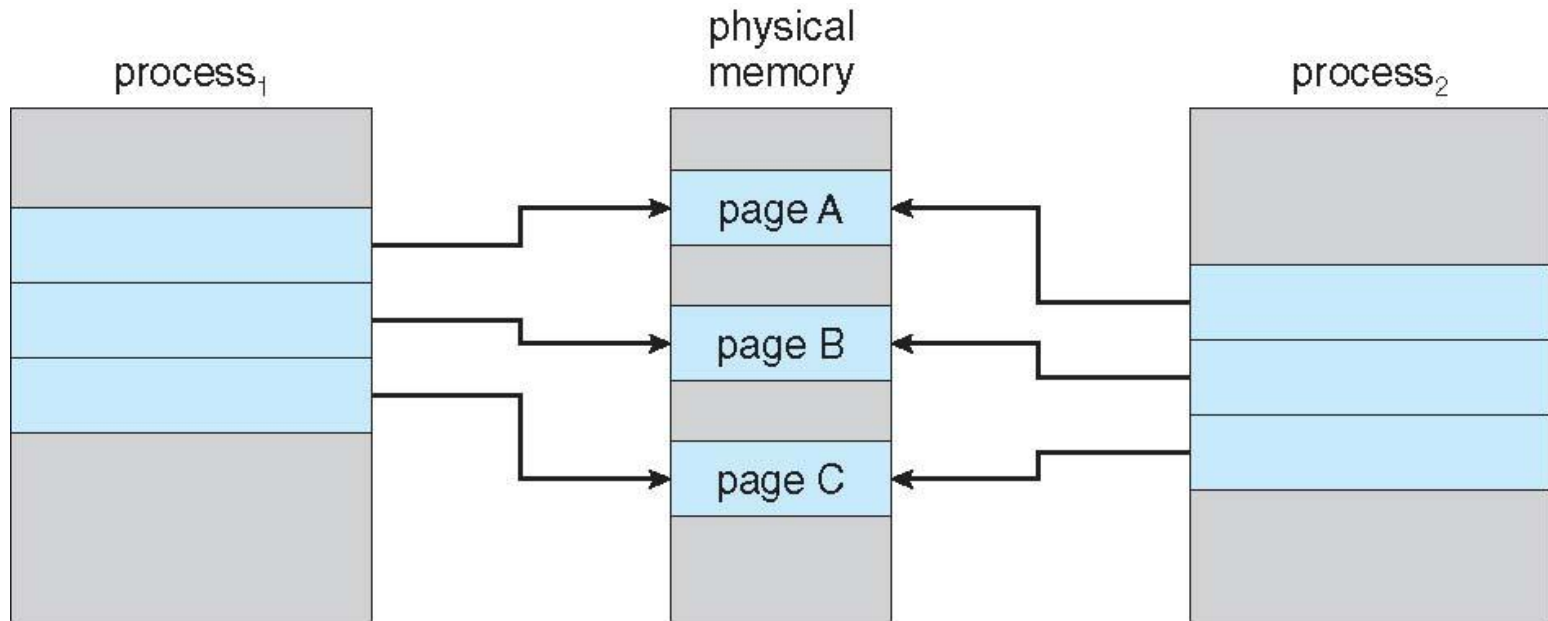
- Extreme case – start process with *no* pages in memory
 - OS sets instruction pointer to first instruction of process, non-memory-resident -> page fault
 - And for every other process pages on first access
 - **Pure demand paging** (처음엔 모두 page fault임)
- Actually, a given instruction could access multiple pages -> multiple page faults
 - Consider fetch and decode of instruction which adds 2 numbers from memory and stores result back to memory
 - Pain decreased because of **locality of reference** (메모리로의 데이터 접근이 지역성을 띠م)
- Hardware support needed for demand paging
 - Page table with valid / invalid bit
 - Secondary memory (swap device with **swap space**)
 - Instruction restart

Copy-on-Write

- **Copy-on-Write** (COW) allows both parent and child processes to initially **share** the same pages in memory
 - If either process modifies a shared page, only then is the page copied
- COW allows more efficient process creation as only modified pages are copied
- In general, free pages are allocated from a **pool** of **zero-fill-on-demand** pages
 - Pool should always have free frames for fast demand page execution
 - ▶ Don't want to have to free a frame as well as other processing on page fault
 - Why zero-out a page before allocating it?
- `vfork()` variation on `fork()` system call has parent suspend and child using copy-on-write address space of parent
 - Designed to have child call `exec()`
 - Very efficient

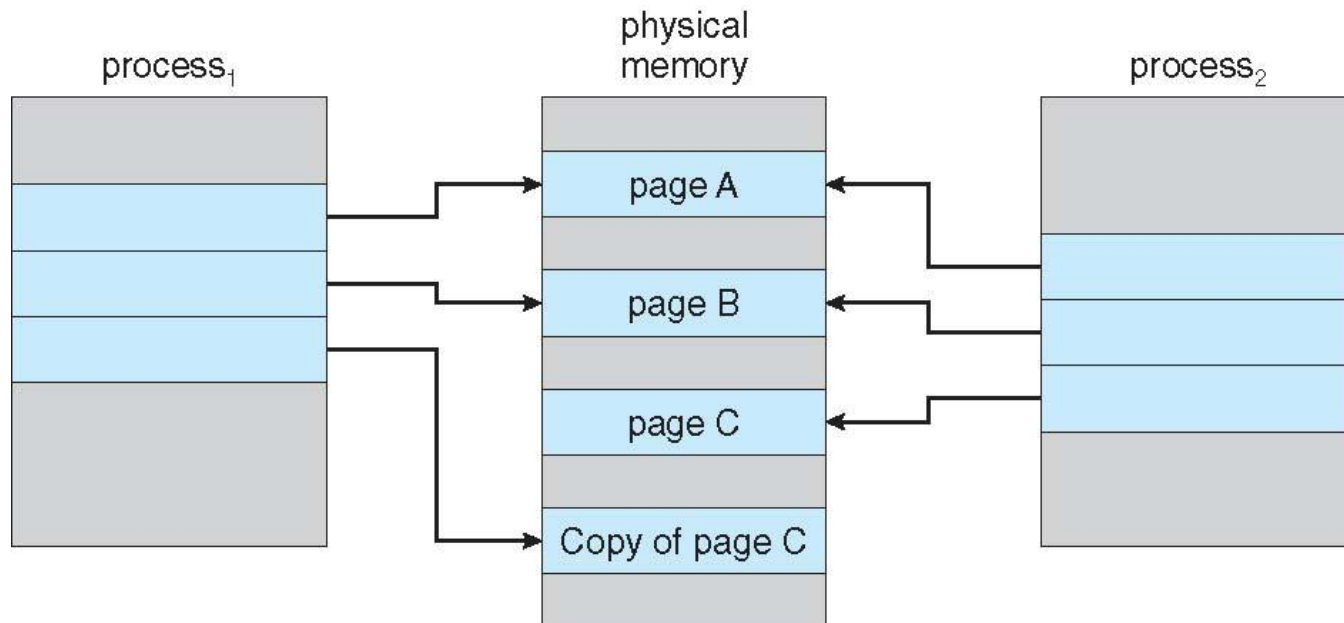
• Q: `fork()`로 생성된 parent/child process들이 page 요구가 많을 때 무언가 효율적인 방법은? > COW

Before Process 1 Modifies Page C



- Q: 메모리를 공유하던 child process가 page 내용을 수정하면?
- > 기본은 read only mode

After Process 1 Modifies Page C



What Happens if There is no Free Frame?


- Used up by process pages
- Also in demand from the kernel, I/O buffers, etc
- How much to allocate to each?
- Page replacement – find some page in memory, but not really in use, page it out
 - Algorithm – terminate? swap out? replace the page?
 - Performance – want an algorithm which will result in minimum number of page faults
- Same page may be brought into memory several times

Page Replacement

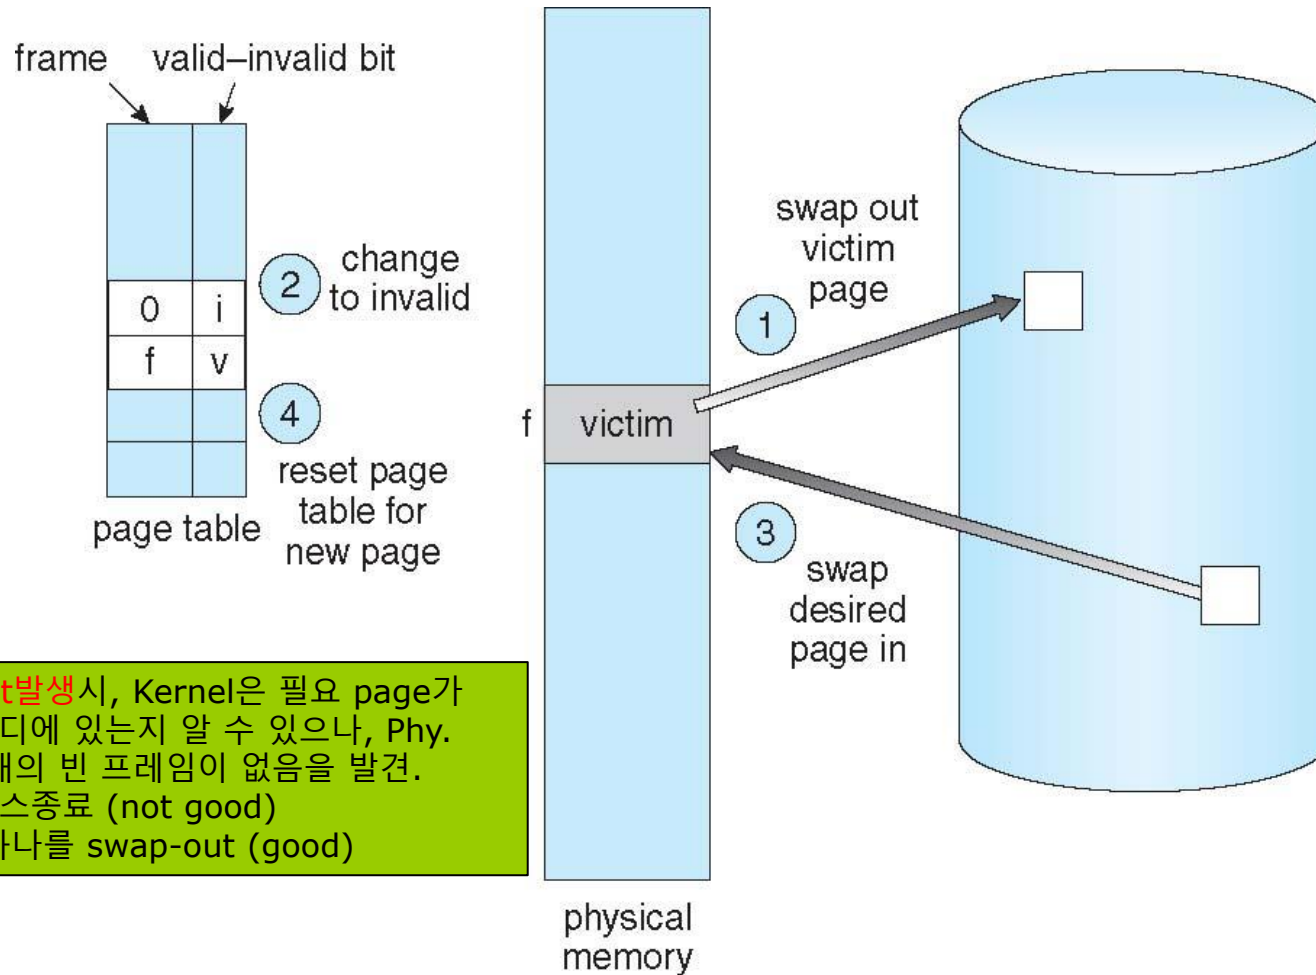
- Prevent **over-allocation** of memory by modifying page-fault service routine to include page replacement
- Use **modify (dirty) bit** to reduce overhead of page transfers – only modified pages are written to disk
- Page replacement completes separation between logical memory and physical memory – large virtual memory can be provided on a smaller physical memory

- Q: Page swap out을 할 때, storage에 write해야 하는데 시간이 걸리는 문제를 해결하려면?
- >dirty bit

Need For Page Replacement



Page Replacement



Page fault발생시, Kernel은 필요 page가 Disk의 어디에 있는지 알 수 있으나, Phy. Memory내의 빈 프레임이 없음을 발견.
 (1)프로세스종료 (not good)
 (2)Page하나를 swap-out (good)

Basic Page Replacement

1. Find the location of the desired page on disk
2. Find a free frame:
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a **victim frame**
 - Write victim frame to disk
3. Bring the desired page into the (newly) free frame; update the page and frame tables
4. Continue the process by restarting the instruction that caused the trap

Note now potentially 2 page transfers for page fault – increasing EAT
(Effective Address Time)

- Q: page swap out/in으로 2번의 copy I/O가 필요하여 성능 문제. 해결은?

Q&A

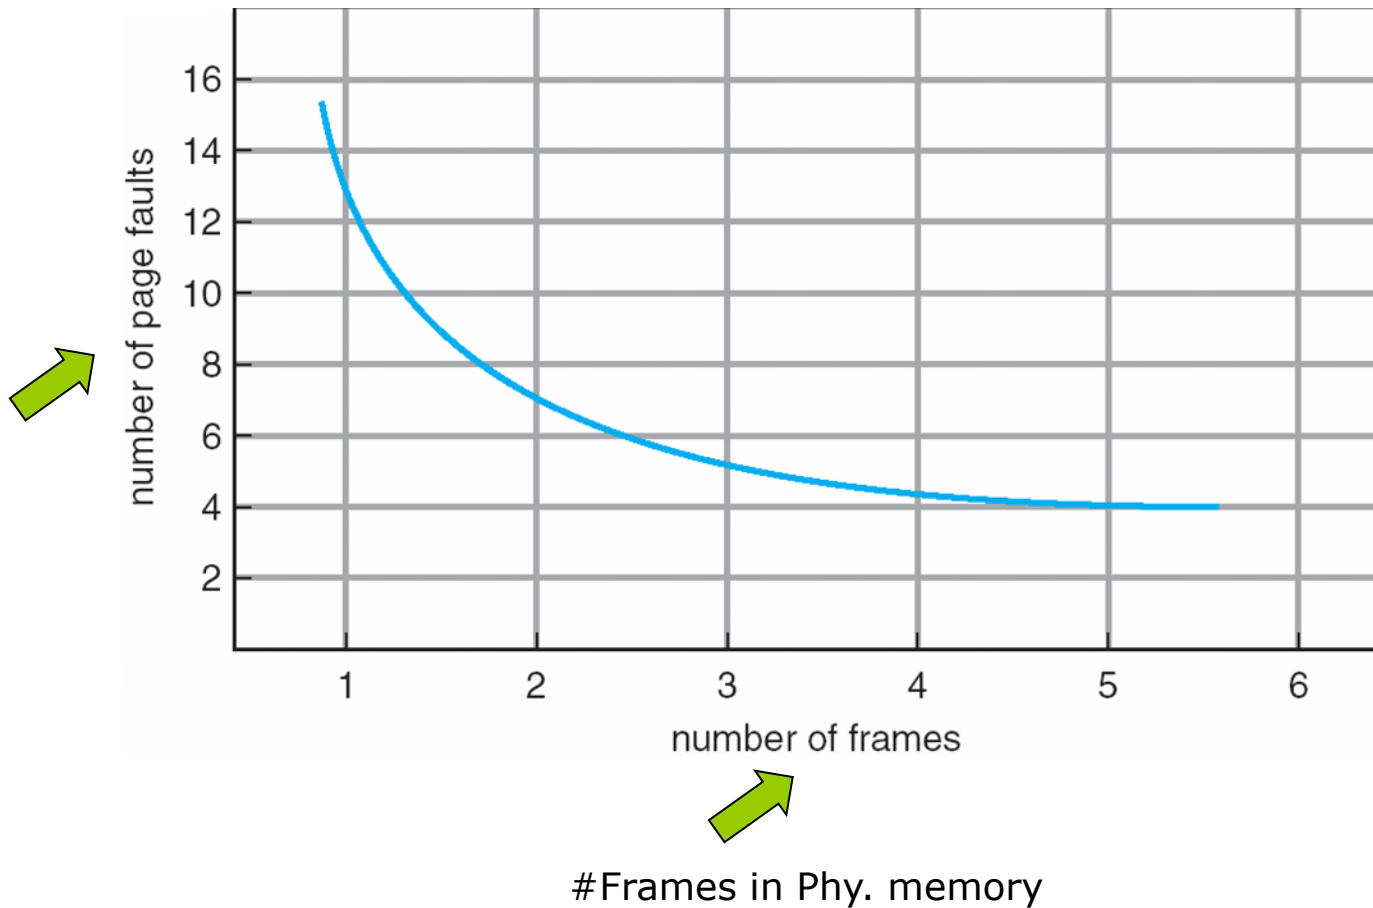
- Question: Page fault를 대응하는 I/O 타임이 너무 긴 문제 – 2번의 I/O필요.
- Ans:
 - CPU가 Page내의 어떤 정보를 update하게 되면, Page 내용이 변경되는 것임. 이를 나타내는 것이 변경비트 (Modify bit or dirty bit) - HW에 존재
 - 만일 Victim Page의 Dirty bit이 true면 메모리의 내용이 원본 디스크상의 내용과 달라짐을 의미. 이 경우 현재 내용을 Disk에 저장할 필요 있음. 하지만, false면 굳이 swap out으로 Disk에 저장할 필요 없음.
 - 따라서, 경우에 따라 Page fault의 I/O시간을 ½로 줄일 수 있음

Page and Frame Replacement Algorithms

- **Frame-allocation algorithm** determines
 - How many frames to give each process
 - Which frames to replace
- **Page-replacement algorithm**
 - Want lowest page-fault rate on both first access and re-access
- Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
 - String is just page numbers, not full addresses
 - Repeated access to the same page does not cause a page fault
 - Results depend on number of frames available
- In all our examples, the **reference string** of referenced page numbers is


7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1

Graph of Page Faults Versus The Number of Frames



First-In-First-Out (FIFO) Algorithm

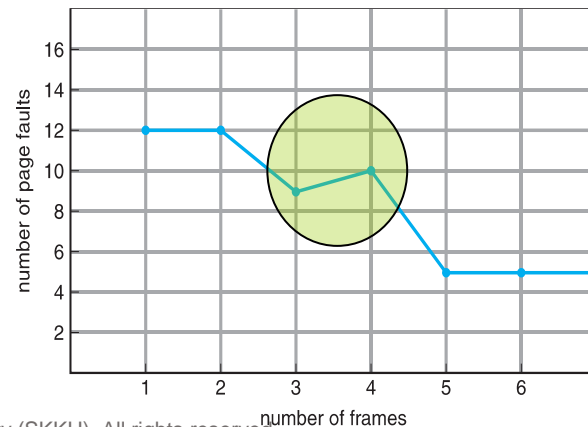
- Reference string: **7,0,1,2,0,3,0,4,2,3,0,3,2,1,2,0,1,7,0,1**
- 3 frames** (3 pages can be in memory at a time per process)



Replace oldest
page slot

15 page faults

- Can vary by reference string: consider 1,2,3,4,1,2,5,1,2,3,4,5
 - Adding more frames can cause more page faults!**
 - Belady's Anomaly**
- How to track ages of pages?
 - Just use a FIFO queue



1,2,3,4,1,2,5,1,2,3,4,5

3 frames

1, 2, 3

4, 2, 3

4, 1, 3

4, 1, 2

5, 1, 2

5, 3, 2

5, 3, 4

9 times

4 frames

1, 2, 3, 4

5, 2, 3, 4

5, 1, 3, 4

5, 1, 2, 4

5, 1, 2, 3

4, 1, 2, 3

4, 5, 2, 3

10 times

더 많은 frame에
도리어 page
fault가 느는 경우

Optimal Algorithm


- Replace page that will not be used for longest period of time
 - 9 is optimal for the example
- How do you know this?
 - Can't read the future
- Used for measuring how well your algorithm performs

앞으로 가장 오랫동안 사용되지 않을 페이지를 찾아 교체

닥터 스트레인지 필요

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1




page frames

Least Recently Used (LRU) Algorithm

- Use past knowledge rather than future
- Replace page that has not been used in the most amount of time
- Associate time of last use with each page

가장 오랫동안 사용되지 않은 페이지를 교체



- 12 faults – better than FIFO but worse than OPT
- Generally good algorithm and frequently used
- But how to implement?

LRU Algorithm (Cont.)

- Counter implementation

- Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
- When a page needs to be changed, look at the counters to find smallest value
 - ▶ Search through table needed


- Stack implementation

- Keep a stack of page numbers in a double link form:
- Page referenced:
 - ▶ move it to the top
 - ▶ requires 6 pointers to be changed
- But each update more expensive
- No search for replacement

LRU는 HW지원 필요 2가지 방법가능
- Counters
- Stack – 페이지 참조될 때 마다
Page 번호가 Stack 중간에서 제거
되어 top에 놓이게 함. ->
Bottom이 가장 오랫동안 안쓴
페이지. : Double linked list로
구현

- LRU and OPT are cases of **stack algorithms** that don't have Belady's Anomaly

Use Of A Stack To Record Most Recent Page References



Counting Algorithms


- Keep a counter of the number of references that have been made to each page
 - Not common 둘다 잘 안쓰임
- **Least Frequently Used (LFU) Algorithm**: replaces page with smallest count
- **Most Frequently Used (MFU) Algorithm**: based on the argument that the page with the smallest count was probably just brought in and has yet to be used

LFU 즉, 참조회수가 가장 작은 페이지를 교체: 활발한 페이지는 참조회수가 많아질 것이란 생각. 하지만, 어떤 프로세스가 초기에 한 페이지 집중적 사용 후 그 후로 사용 잘 안하면 문제. 계속 메모리에 머물게 됨.
해결책 > 참조회수를 일정시간마다 right shift하여 지수적으로 영향력 감소시킴 (일종의 aging의 HW사용한 변형)

Thrashing

- If a process does not have “enough” pages, the page-fault rate is very high

- Page fault to get page
- Replace existing frame
- But quickly need replaced frame back
- This leads to:
 - ▶ Low CPU utilization
 - ▶ Operating system thinking that it needs to increase the degree of multiprogramming
 - ▶ Another process added to the system



- **Thrashing** \equiv a process is busy swapping pages in and out

요약(1/2)

- 가상메모리는 프로세스 전체가 메모리 내에 올라오지 않더라도 실행이 가능하도록 하는 기법
 - 사용자 프로그램이 물리 메모리 보다 커도 됨
- 각 프로세스들은 공유 라이브러리가 자신의 일부라고 생각하지만, 실제로는 라이브러리가 존재하는 물리 메모리 페이지들은 모든 프로세스들에게 공유 (라이브러리는 읽기만이 허용되는 상태)
- 요구페이징(Demand Paging) – Swapping 기법과 유사. 통상 프로세스는 보조 메모리(디스크)에 존재하나, 프로세스를 실행하고 싶으면 메모리로 읽어들임(Swap in). 단, 필요한 Page 만 읽어들임.
- Page Table의 유효비트/무효비트는 해당페이지가 메모리에 있다 없음을 의미.
- 메모리에 올라와 있지 않는 페이지에 프로세스가 접근하면, Page fault trap 발생 -> Paging HW가 주소변환 과정에서 무효 비트 발견하고 OS에 trap. 유효한 참조인데 아직 Page가 메모리에 없다면 디스크로 부터 가져옴.

요약(2/2)

- Locality of Reference: 통상 프로세스들을 분석하면, 프로그램의 어느 한 특정 작은 부분만 한동안 집중적으로 참조 -> Demand Paging (필요한 부분만 Paging)은 좋은 성과
- Copy-on-Write: 예로 부모/자식 프로세스처럼 fork() 때 주소공간을 초기에 복사하는 경우, 자식들은 부모로부터 Page들을 복사할 필요 없이 share해서 쓰다가 data를 write해서 메모리가 변경이 되는 때에만 Copy-on-write로 해당 페이지를 복사하여 write. : vfork() 등이 존재
- Page Replacement: 만약 빈 프레임이 없다면, 현재 사용되지 않는 프레임을 찾아서 그것을 비움. 그 프레임의 내용을 swap공간에 쓰고, 그 페이지가 메모리에 더 이상 존재하지 않음을 Page table 수정하여 프레임을 빈 상태로 만듦. 그 후 빈 프레임에 page fault난 page를 저장
- Page fault의 I/O를 줄이는 법: HW기반의 Dirty-bit사용 (변경 유무 확인)